

GuessTuples Project

Andrew J. Wren

5 May 2021

Abstract

Notes on GuessTuples project aka NLearn

1 Configuring the nets

1.1 Alice

One per bit. The input array to guess is $\mathbf{x} = (x_j)_{j=0,\dots,N_{\text{elements}}}$. There should be N_{code} outputs taking values $\mathbf{y} = (y_j)_{j=0,\dots,N_{\text{code}}}$.

Normalise all the rewards so that for each bit j , $r_{j\checkmark} + (N_{\text{code}} - 1) r_{j\text{x}} = 0$. In other words

$$r_{jk} \leftarrow r_{jk} - \frac{r_{j\checkmark} + (N_{\text{code}} - 1) r_{j\text{x}}}{N_{\text{code}}}. \quad (1)$$

The Q estimate is then taken to be

$$Q(\mathbf{x}) = \sum_j b_j y_j \equiv \sum_j |y_j|, \quad (2)$$

where

$$b_j = \text{sgn}(y_j) \quad (3)$$

is the prediction for the machine value of the j th bit. The loss function is

$$L = |Q(\mathbf{x}) - r|^2. \quad (4)$$

Alternative approaches include:

1. Two outputs for each bit showing the reward for each of 0 and 1. *May reflect negative rewards better?*
2. Combine the rewards from the bits (with either one or two outputs per bit) by something other than addition - e.g. multiplication or via an NN. *The NN option seems quite interesting. Interesting to use pytorch's gradients for that.*
3. **One per code.** One output for each possible code. *Might work but $2^{N_{\text{code}}}$ is quite large... not impossibly so if $N_{\text{code}} = 8$.*

4. Inspired by Ref. [1], feed \mathbf{x} into Alice's 'first' net, to get output \mathbf{y} , and all possible codes \mathbf{c} into her 'second' net, both net's having the same target dimensionality (a hyperparameter). Then the code to use is the one $\mathbf{c}(\mathbf{x})$ closest to the output of the first net, with the Q being given by the inner product $Q = \langle \mathbf{y}, \mathbf{c}(\mathbf{x}) \rangle$. Ref. [2] might provide an alternative, actor-critic, approach on a similar theme. The main case above is, in effect, an embedding of \mathbf{x} into the target space (of dimensionality N_{code}) which then compares with the natural embedding of \mathbf{c} by, in effect, the inner product.
5. Ref. [3] suggest sequentialising, which points to a variant of our main approach which does each bit in succession and feeding those results into successive Alice-nets so the Q -estimate for later bits takes account of earlier bits / estimates, with the N_{code} th estimate providing a final code \mathbf{c} and Q -estimate for that code.
6. Move away from typical Q-learning. Instead Alice's output is the code \mathbf{c} and then when Bob makes his choice \mathbf{x}_{pred} (see below) run that choice through a copy of Alice, to get \mathbf{c}_{Bob} and then the loss function is

$$L = -r(\mathbf{c}, \mathbf{c}_{\text{Bob}}). \quad (5)$$

1.2 Bob

One per bit aka Simple. Bob receives a matrix, $\mathbf{X} = (X_i) = (X_{ij})$ for $0 \leq i < N_{\text{select}}$, $0 \leq j < N_{\text{elements}}$, and a code $\mathbf{c} = (c_k)_{k=0, \dots, N_{\text{code}}}$. Why not makes his outputs be Q -estimates $\mathbf{z} = (z_i)_{i=0, \dots, N_{\text{select}}}$. Bob's prediction is then $\mathbf{x}_{\text{pred}} = \mathbf{X}_{i_{\text{pred}}}$ where

$$i_{\text{pred}} = \text{argmax}_i (z_i). \quad (6)$$

The loss function is

$$L = |z_{i_{\text{pred}}} - r|^2. \quad (7)$$

How do we enforce covariance with respect to the order of (X_i) ?

1. Covariance will occur naturally and quickly without any specific intervention. *To be determined.*
2. Covariance can be enforced through choosing a set $\{\sigma\} \subseteq S_{n_{\text{code}}}$, which could be generated element-by-element by composing randomly-selected basis transpositions $(j \ j+1)$, and then adding to the loss a term

$$\mu \sum_{\sigma} |z - \sigma^{-1} [z(\sigma[\mathbf{X}])]|^2 \quad (8)$$

for some fixed hyperparameter $\mu > 0$. Note this the term is still run backward through the original $\mathbf{x} \mapsto \mathbf{z}$ net configuration only. *How effective would that be? How big does $\{\sigma\}$ have to be? And how much time would the permutation and the additions forward passes cost?*

3. Enforce covariance via direct identification of weights in Bob's net. *How?*
4. Something related to set transformers. ?
5. Adopt a different basic set-up where each (X_i) is fed through the net separately, alongside the code \mathbf{c} , resulting in a Q -estimate z_i . Then find the loss function as in Eq. 7. *Seems the most straightforward?*

None of these quite amount to Bob seeks to reproduce the Alice's code vocabulary. However Bob could additionally set up a net in the same basic configuration as Alice's (he doesn't know the weights of course) and train *that* net jointly with his main net.




Figure 1: The best results — from /runs/Apr27_23-01-58_andrew-XPS-15-9570. The lines show the square root of the mean square losses with (a) $\text{lr}=0.3$ Alice (orange), Bob (dark blue); (b) $\text{lr}=0.1$ Alice (brick red), Bob (cyan); (c) $\text{lr}=0.01$ Alice (pink), Bob (green). The plot is from TensorFlow and uses smoothing of 0.999. Note rewards from random plays are counted.

2 Results

2.1 Original strategies

Figure 1 is representative of the better results for the original strategies, **one per bit** — in other words, not very good.¹ Increasing from `h.GAMESIZE = 1` to `h.GAMESIZE = 32` gives no better results.

3 Revised approach — NLearn

Key runs:

1. 21-05-01_12:05:16 is the strategy that works

```
'ALICE_STRATEGY': 'from_decisions',
'BOB_STRATEGY': 'circular_vocab'
```

up to a point when it levels off. Gets to reward = 0.6.

2. 21-05-01_20:04:35 other lr choices but same result — see Figure 2
3. 21-05-02_17:29:40 stops Alice training at some point. Alice $\text{lr} = 0.1$ and Bob $\text{lr} = 0.01$ gets to 0.8 — see Figure 3. The config includes

```
hyperparameters = {
```

¹The plot is taken from TensorBoard which gives an .svg file, then converted to .pdf by `rsvg-convert -f pdf -o <fig-file-name>.pdf "Sqrt losses.svg"`.



Figure 2: Mean Rewards per game for 21-05-01_20:04:35. By colour, (Alice \backslash r, Bob \backslash r) are: cyan (0.1, 0.1), orange (0.1, 0.01), pink (0.01, 0.1), and blue (0.01, 0.01). Note rewards from random plays are counted.



Figure 3: Mean Rewards per game for 21-05-02_17:29:40. By colour, (Alice \backslash r, Bob \backslash r) are: green (0.1, 0.1), orange (0.1, 0.01), grey (0.01, 0.1), and cyan (0.01, 0.01). Note rewards from random plays are counted.

```

'N_ITERATIONS': 500000,
'RANDOM_SEED': 42,
'TORCH_RANDOM_SEED': 4242,
'ALICE_LAYERS': 3,
'ALICE_WIDTH': 50,
'BOB_LAYERS': 3,
'BOB_WIDTH': 50,
'BATCHSIZE': 32,
'GAMESIZE': 32,
'BUFFER_CAPACITY': 640000,
'START_TRAINING': 20000,
'N_SELECT': 5,
'EPSILON_ONE_END': 40000,
'EPSILON_MIN': 0.01,
'EPSILON_MIN_POINT': 300000,
'ALICE_STRATEGY': 'from_decisions',
'BOB_STRATEGY': 'circular_vocab',
'ALICE_OPTIMIZER': ('SGD', '{"lr": 0.1}'),
'BOB_OPTIMIZER': ('SGD', '{"lr": 0.01}'),
'ALICE_LOSS_FUNCTION': ('MSE', {}),
'BOB_LOSS_FUNCTION': 'Same',
'ALICE_LAST_TRAINING': 200000

```

Alice here, 21-05-02_17:29:40 hp_run=2 generates codes as follows:

```

11101100 [0, 1, 2, 12, 13, 14, 15]
11101110 [3]
10101110 [4, 6]
10100110 [5, 7]
10100100 [8]
11100100 [9, 10, 11]

```

Surprisingly only six distinct codes used! At least the first and last have sequential runs of numbers.

4. If increase N_SELECT to 16 (all the numbers shown to Bob), then, in run 21-05-03_10:53:10, gets to reward = 0.8, as good as for N_SELECT = 5. In fact very slightly better (mean at 25.0° rather than 32.9°) Alice's code book is still very small:

```

21-05-03_10:53:10BST_NLearn_model_1_Alice_iter500000

01111010 [0, 15]
01111100 [1, 2, 3, 4, 5, 14]
01011100 [6, 7]
01011110 [8, 9, 12, 13]
01111110 [10, 11]

```



Figure 4: With `N_SELECT = 16`, at 21-05-03_10:53:10.

4 From now on exclude random plays from mean reward

The exclusion is if either Alice or Bob or both is random.

4.1 Loss includes element to push bits towards -1 or 1 , and simple ‘proximity bonus’

This gets pretty good results — see Figure 5 which also (orange, pink, blue) lines adds a ‘proximity bonus’ that — at least for these seeds — speeds up training but does not improve the outcome.

4.2 Loss includes element to push bits towards -1 or 1 , and simple ‘proximity bonus’

At 21-05-05_11:27:12, changing Section 4.1 by

```
'N_ITERATIONS': 15 * (10 ** 4),
'ALICE_PROXIMITY_BONUS': 30000,
'ALICE_PROXIMITY_SLOPE_LENGTH': 10 ** 4
```

get the excellent result shown in Figure 6, having a final smoothed value of 0.94. The final coding and decoding books are

```
00100111 [0, 1, 2, 3]
10100111 [4]
10110111 [5]
10110011 [6, 7]
10111011 [8]
10101010 [9, 10]
10101101 [11]
```




Figure 5: The green line shows the best run from 21-05-03_20:36:57, which introduced MSEBits and had Alice stopping training at iteration 300 000. The remaining lines are from 21-05-04_20:10:38 and do not stop Alice training. They add the simple ‘proximity bonus’ of 1 when codes or numbers are equal — phasing this in over 10 000 iterations from iteration 100 000 (orange), 200 000 (pink) and 300 000 (blue), The plot has smoothing set to 0.9.

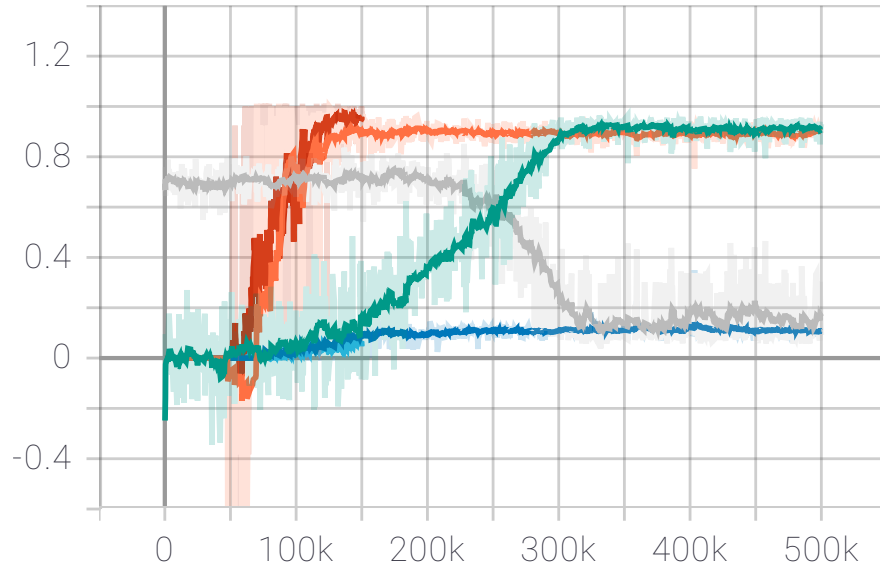


Figure 6: The red line shows the mean reward of 21-05-05_11:27:12, while the just visible cyan line is its standard deviation. The orange and green lines are as in Figure 5, with the blue and grey lines being their respective standard deviations. The plot has smoothing set to 0.9.

```
10100101 [12]
00100101 [13, 14, 15]
```

```
00100111 2
10100111 4
10110111 5
10110011 6
10111011 8
10101010 10
10101101 11
10100101 12
00100101 14
```

with Alice using nine codes.

5 Things to try

1. What codes does best Alice generate?
2. Try using the loss function to constraint outputs to nearer bit values. Try increasing the weighting of this.
3. How quickly can `epsilon` be tapered?
4. Vary learning rates.
5. Vary `modulus`, `N_CODE` and `N_SELECT`.
6. Introduce noise.
7. Alice strategy with a code, as input and the output are values for the numbers. In each play (or train?) step feed all the codes in and the outputs indicate how well represents each number???
8. Try best strategy but with Alice outputs having dimension $2 \times N_CODE$.
9. Train bits successively.
10. Look at MARL literature.
11. (At some stage in the training) introduce a ‘proximity bonus’ into Alice’s training, which increases (in the same way) both the closeness of codes and the rewards if Bob’s decision is right or nearly so.
12. Do a second sweep of `epsilon` going from high to low — perhaps for one player only? Definitely should re-`epsilon`-randomise Bob as otherwise Alice will never (or rarely if `N_SELECT < N_CODE`) get fed choices not in his decoding book. And I think Alice too, so Bob can learn new codes.
13. Random seeds seem to play a significant role — at least for short (~ 12 500) iteration training. Test how significant for 500 000 iterations.

References

- [1] J. He, J. Chen, X. He, J. Gao, L. Li, L. Deng et al., *Deep reinforcement learning with a natural language action space*, *arXiv preprint arXiv:1511.04636* (2015) .
- [2] G. Dulac-Arnold, R. Evans, H. van Hasselt, P. Sunehag, T. Lillicrap, J. Hunt et al., *Deep reinforcement learning in large discrete action spaces*, *arXiv preprint arXiv:1512.07679* (2015) .
- [3] S. J. Majeed and M. Hutter, *Exact reduction of huge action spaces in general reinforcement learning*, *arXiv preprint arXiv:2012.10200* (2020) .