GuessTuples Project

Andrew J. Wren

15 May 2021

Abstract

Notes on `GuessTuples` project aka `NLearn`

1 Configuring the nets

1.1 Alice

One per bit. The input array to guess is $x = (x_j)_{j=0,\dots,N_{\text{elements}}}$. There should be N_{code} outputs taking values $y = (y_j)_{j=0,\dots,N_{\text{code}}}$.

Normalise all the rewards so that for each bit j, $r_{j\checkmark} + (N_{\text{code}} - 1)\, r_{j\times} = 0$. In other words

$$r_{jk} \leftarrow r_{jk} - \frac{r_{j\checkmark} + (N_{\text{code}} - 1)\, r_{j\times}}{N_{\text{code}}}. \tag{1}$$

The Q estimate is then taken to be

$$Q(x) = \sum_j b_j y_j \equiv \sum_j |y_j|, \tag{2}$$

where

$$b_j = \text{sgn}(y_j) \tag{3}$$

is the prediction for the machine value of the jth bit. The loss function is

$$L = |Q(x) - r|^2. \tag{4}$$

Alternative approaches include:

1. Two outputs for each bit showing the reward for each of 0 and 1. *May reflect negative rewards better?*

2. Combine the rewards from the bits (with either one or two outputs per bit) by something other than addition - e.g. multiplication or via an NN. *The NN option seems quite interesting. Interesting to use `pytorch`'s gradients for that.*

3. **One per code.** One output for each possible code. *Might work but $2^{N_{code}}$ is quite large... not impossibly so if $N_{code} = 8$.*

4. Inspired by Ref. [1], feed x into Alice's 'first' net, to get output y, and all possible codes c into her 'second' net, both net's having the same target dimensionality (a hyperparameter). Then the code to use is the one $c(x)$ closest to the output of the first net, with the Q being given by the inner product $Q = \langle y, c(x) \rangle$. Ref. [2] might provide an alternative, actor–critic, approach on a similar theme. The main case above is, in effect, an embedding of x into the target space (of dimensionality N_{code}) which then compares with the natural embedding of c by, in effect, the inner product.

5. Ref. [3] suggest sequentialising, which points to a variant of our main approach which does each bit in succession and feeding those results into successive Alice–nets so the Q-estimate for later bits takes account of earlier bits / estimates, with the N_{code}th estimate providing a final code c and Q-estimate for that code.

6. Move away from typical Q-learning. Instead Alice's output is the code c and then when Bob makes his choice x_{pred} (see below) run that choice through a copy of Alice, to get c_{Bob} and then the loss function is

$$L = -r(c, c_{\text{Bob}}). \tag{5}$$

1.2 Bob

One per bit aka **Simple.** Bob receives a matrix, $\mathbf{X} = (X_i) = (X_{ij})$ for $0 \le i < N_{\text{select}}$, $0 \le j < N_{\text{elements}}$, and a code $c = (c_k)_{k=0,\dots,N_{\text{code}}}$. Why not makes his outputs be Q-estimates $z = (z_i)_{i=0,\dots,N_{\text{select}}}$. Bob's prediction is then $x_{\text{pred}} = X_{i_{\text{pred}}}$ where

$$i_{\text{pred}} = \text{argmax}_i(z_i). \tag{6}$$

The loss function is

$$L = \left| z_{i_{\text{pred}}} - r \right|^2. \tag{7}$$

How do we enforce covariance with respect to the order of (X_i)?

1. Covariance will occur naturally and quickly without any specific intervention. *To be determined.*

2. Covariance can be enforced through choosing a set $\{\sigma\} \subseteq S_{n_{\text{code}}}$, which could be generated element–by–element by composing randomly–selected basis transpositions $(j \ j+1)$, and then adding to the loss a term

$$\mu \sum_\sigma \left| z - \sigma^{-1} \left[z(\sigma[\mathbf{X}]) \right] \right|^2 \tag{8}$$

for some fixed hyperparameter $\mu > 0$. Note this the term is still run backward through the original $x \mapsto z$ net configuration only. *How effective would that be? How big does $\{\sigma\}$ have to be? And how much time would the permutation and the additions forward passes cost?*

3. Enforce covariance via direct identification of weights in Bob's net. *How?*

4. Something related to set transformers. *?*

5. Adopt a different basic set–up where each (X_i) is fed through the net separately, alongside the code c, resulting in a Q-estimate z_i. Then find the loss function as in Eq. 7. *Seems the most straightforward?*

None of these quite amount to Bob seeks to reproduce the Alice's code vocabulary. However Bob could additionally set up a net in the same basic configuration as Alice's (he doesn't know the weights of course) and train *that* net jointly with his main net.

Figure 1: The best results — from `/runs/Apr27_23-01-58_andrew-XPS-15-9570`. The lines show the square root of the mean square losses with (a) `lr=0.3` Alice (orange), Bob (dark blue); (b) `lr=0.1` Alice (brick red), Bob (cyan); (c) `lr=0.01` Alice (pink), Bob (green). The plot is from TensorFlow and uses smoothing of 0.999. Note rewards from random plays are counted.

2 Results

2.1 Original strategies

Figure 1 is representative of the better results for the original strategies, **one per bit** — in other words, not very good.[1] Increasing from `h.GAMESIZE = 1` to `h.GAMESIZE = 32` gives no better results.

3 Revised approach — `NLearn`

Key runs:

1. `21-05-01_12:05:16` is the strategy that works

   ```
   'ALICE_STRATEGY': 'from_decisions',
   'BOB_STRATEGY': 'circular_vocab'
   ```

   up to a point when it levels off. Gets to `reward = 0.6`.

2. `21-05-01_20:04:35` other `lr` choices but same result — see Figure 2

3. `21-05-02_17:29:40` stops Alice training at some point. Alice `lr = 0.1` and Bob `lr = 0.01` gets to 0.8 — see Figure 3. The `config` includes

   ```
   hyperparameters = {
   ```

---

[1]The plot is taken from TensorBoard which gives an `.svg` file, then converted to `.pdf` by `rsvg-convert -f pdf -o <fig-file-name>.pdf "Sqrt losses.svg"`.

Figure 2: Mean Rewards per game for `21-05-01_20:04:35`. By colour, (Alice `lr`, Bob `lr`) are: cyan $(0.1, 0.1)$, orange $(0.1, 0.01)$, pink $(0.01, 0.1)$, and blue $(0.01, 0.01)$. Note rewards from random plays are counted.

Figure 3: Mean Rewards per game for `21-05-02_17:29:40`. By colour, (Alice `lr`, Bob `lr`) are: green $(0.1, 0.1)$, orange $(0.1, 0.01)$, grey $(0.01, 0.1)$, and cyan $(0.01, 0.01)$. Note rewards from random plays are counted.

```
'N_ITERATIONS': 500000,
'RANDOM_SEED': 42,
'TORCH_RANDOM_SEED': 4242,
'ALICE_LAYERS': 3,
'ALICE_WIDTH': 50,
'BOB_LAYERS': 3,
'BOB_WIDTH': 50,
'BATCHSIZE': 32,
'GAMESIZE': 32,
'BUFFER_CAPACITY': 640000,
'START_TRAINING': 20000,
'N_SELECT': 5,
'EPSILON_ONE_END': 40000,
'EPSILON_MIN': 0.01,
'EPSILON_MIN_POINT': 300000,
'ALICE_STRATEGY': 'from_decisions',
'BOB_STRATEGY': 'circular_vocab',
'ALICE_OPTIMIZER': ('SGD', '{"lr": 0.1}'),
'BOB_OPTIMIZER': ('SGD', '{"lr": 0.01}'),
'ALICE_LOSS_FUNCTION': ('MSE', {}),
'BOB_LOSS_FUNCTION': 'Same',
'ALICE_LAST_TRAINING': 200000
```

Alice here, `21-05-02_17:29:40 hp_run=2` generates codes as follows:

```
11101100 [0, 1, 2, 12, 13, 14, 15]
11101110 [3]
10101110 [4, 6]
10100110 [5, 7]
10100100 [8]
11100100 [9, 10, 11]
```

Surprisingly only six distinct codes used! At least the first and last have sequential runs of numbers.

4. If increase `N_SELECT` to 16 (all the numbers shown to Bob), then, in run `21-05-03_10:53:10`, gets to reward = 0.8, as good as for `N_SELECT` = 5. In fact very slightly better (mean at 25.0° rather than 32.9°) Alice's code book is still very small:

```
21-05-03_10:53:10BST_NLearn_model_1_Alice_iter500000

01111010 [0, 15]
01111100 [1, 2, 3, 4, 5, 14]
01011100 [6, 7]
01011110 [8, 9, 12, 13]
01111110 [10, 11]
```

Figure 4: With `N_SELECT` = 16, at `21-05-03_10:53:10`.

4 From now on exclude random plays from mean reward

The exclusion is if either Alice or Bob or both is random.

4.1 Loss includes element to push bits towards −1 or 1, and simple 'proximity bonus'

This gets pretty good results — see Figure 5 which also (orange, pink, blue) lines adds a 'proximity bonus' that — at least for these seeds — speeds up training but does not improve the outcome.

4.2 Loss includes element to push bits towards −1 or 1, and simple 'proximity bonus'

At `21-05-05_11:27:12`, changing Section 4.1 by

```
'N_ITERATIONS': 15 * (10 ** 4),
'ALICE_PROXIMITY_BONUS': 30000,
'ALICE_PROXIMITY_SLOPE_LENGTH': 10 ** 4
```

get the excellent result shown in Figure 6, having a final smoothed value of 0.94. The final coding and decoding books are

```
00100111 [0, 1, 2, 3]
10100111 [4]
10110111 [5]
10110011 [6, 7]
10111011 [8]
10101010 [9, 10]
10101101 [11]
```

Figure 5: The green line shows the best run from `21-05-03_20:36:57`, which introduced `MSEBits` and had Alice stopping training at iteration 300 000. The remaining lines are from `21-05-04_20:10:38` and do not stop Alice training. They add the simple 'proximity bonus' of 1 when codes or numbers are equal from iteration 100 000 (orange), 200 000 (pink) and 300 000 (blue), The plot has smoothing set to 0.9.

Figure 6: The red line shows the mean reward of `21-05-05_11:27:12`, while the just visible cyan line is its standard deviation. The orange and green lines are as in Figure 5, with the blue and grey lines being their respective standard deviations. The plot has smoothing set to 0.9.

```
10100101 [12]
00100101 [13, 14, 15]


00100111 2
10100111 4
10110111 5
10110011 6
10111011 8
10101010 10
10101101 11
10100101 12
00100101 14
```

with Alice using nine codes.

However, another run, `21-05-05_13:13:06`, with the same parameters, except for the three seeds, shows the high random dependence getting a small code book:

```
10010100 [0, 1, 2, 3, 12, 13, 14, 15]
10010000 [4, 5, 6, 7]
00111000 [8, 9, 10]
10110100 [11]


10010100 0
10010000 5
00111000 9
10110100 11
```

Figure 7 compares with previous results. Perhaps suggests introducing some noise?

4.3   Noise

From `21-05-05_21:56:16`, noise doesn't seem to help on this individual run — see Figure 8. However, does it make the model more robust to changes in random seeds?

4.4   In Alice training, make both sides of the loss function have `grad`

As in `21-05-06_09:44:04`, this doesn't work — reward oscillates around zero. Is also slower.

4.5   Phasing in proximity bonus, double deep learning for Alice

In `21-05-06_20:41:41` find that phasing in (over 10 000 iterations) helps — getting to 0.96 — but adding double deep learning for Alice may not. See Figure 9.

Figure 7: From `21-05-05_13:13:06` we have the red line. The grey line which is the former run shown in red in Figure 6 and the that shown in orange in Figure 6. The plot has smoothing set to 0.9.


Figure 8: From `21-05-05_21:56:16`, with noise of 0.01 (green), 0.03 (orange), 0.1 (grey), starting at iteration 175 000 (earlier starts were poorer). The noise is cut off before the end to allow comparison. Plot smoothing is at 0.9, and lesser smoothing doesn't show more post–noise recovery.

Figure 9: From `21-05-06_20:41:41`, all have phased–in proximity bonus with `ALICE_DOUBLE` set to None (green), 1000 (pink), 100 (blue) and 5000 (orange). Smoothing 0.9.

5   After correcting error in the construction of `MSEBits`

Just before `21-05-07_14:03:47`, corrected an error in the construction of `MSEBits` which meant it was looking at `closeness` not `alice_outputs_from_target` for the bits! This meant that trying to push closeness towards ±1 whatever it should have been! Surprised it was so successful *before*. However, `21-05-07_14:03:47` is very unsuccessful, mean reward just oscillating around zero — this is because I messed up taking the means in commit `a3038c3` and I think previously.

That previous loss function was

```
self.loss_fn = lambda x, y: torch.mean(
    torch.sum(
        torch.nn.functional.mse_loss(x, y, reduction='none')
        + (mu / 2) * torch.square(x - torch.sign(x)),
        dim=-1
    )
)
```

which had the effect of trying to push the `closeness` to ±1 with the second term, and — taking a sum over batches to be followed by a mean over a scalar — multiplying the learning rate by 32.

The 'correct' formulation is

```
self.loss_fn = (
    lambda x, y, z:
    torch.nn.functional.mse_loss(x, y)
    + (mu / 2)
    * torch.nn.functional.mse_loss(z, torch.sign(z))
) # for both the mse_loss functions this implies reduction='mean'
```

Figure 10: The rewards for `21-05-08_15:50:47`, for the three hp runs: 1 (pink), 2 (blue) and 3 (green) with different tuples of random seeds. Smoothing 0.9.

with z representing the raw output.

Reverting to the 'previous' get a pretty good run at `21-05-07_17:41:29`, but now have code and decode books on view in the run — and Alice's code book doesn't change at all!!!

6   QPerCode

`QPerCode` is a strategy for Alice with its net outputting $2^{N_{code}}$ Q–values. For SGD learning rates of 0.1 and above doesn't do well, but with $lr = 0.001$ at `21-05-08_14:45:11BST` (and double with period 500 and Huber `beta=0.5`) get good results, peaking at reward of 0.870.

The run `21-05-08_15:50:47` for three choices of random seeds shows that this approach is robust. It's also quick (50 000 iteration). See Figure 10. The config has

```
hyperparameters = {
    'N_ITERATIONS': 50 * (10 ** 3), # 5 * (10 ** 5),
    'RANDOM_SEEDS': [
    (868291, 274344, 358840, 94453),
    (382832, 68444, 754888, 857796),
    (736520, 815195, 305871, 974216)
    ],
    'ALICE_NET': 'FFs(3, 50)',
    'BOB_LAYERS': 3,
    'BOB_WIDTH': 50,
    'BATCHSIZE': 32,
    'GAMESIZE': 32,
    'BUFFER_CAPACITY': 32 * 20000,
    'START_TRAINING': 20000,
    'N_SELECT': 16,
```

```
    'EPSILON_ONE_END': 2000,
    'EPSILON_MIN': 0.01,
    'EPSILON_MIN_POINT': 40000,
    'ALICE_PLAY': 'QPerCode',
    'ALICE_TRAIN': 'QPerCode',
    'BOB_STRATEGY': 'circular_vocab',
    'ALICE_OPTIMIZER': [
    'SGD(lr=0.01)'
    ],
    'BOB_OPTIMIZER': [
    ('SGD', '{"lr": 0.01}')
    ],
    'ALICE_LOSS_FUNCTION': 'Huber(beta=0.5)',
    'BOB_LOSS_FUNCTION': ('torch.nn.MSE', {}),
    'ALICE_PROXIMITY_BONUS': 10 ** 8,
    'ALICE_PROXIMITY_SLOPE_LENGTH': 10000,
    'ALICE_LAST_TRAINING': 100 * (10 ** 5),
    'NOISE_START': 10 ** 8,
    'NOISE': 0.,
    'ALICE_DOUBLE': 500
  }

  TUPLE_SPEC = (
  (16,),
  )
  N_CODE = 8

  SMOOTHING_LENGTH = 10000
  SAVE_PERIOD = 10 ** 5
```

Run `21-05-08_17:18:16` for the first tuple of random seeds, varying the Huber `beta` as in Figure 11 suggests `beta=0.1` may be best.

6.1 The first ever 1!

On run `21-05-08_18:30:55` (with a new tuple of seeds) using noise starting at `30000` gives for the first time a reward of close to 1 with Alice's code book using 16 codes for the 16 numbers for the three with non–zero noise (of 0.1, 0,2 and 0.3). Gets

```
   ---- Table of results ----

 code hp_run noise result
 00000   1     0.0 (-0.901, 50000)
 00001   2     0.1 (-0.990, 70000)
 00002   3     0.2 (-0.972, 70000)
 00003   4     0.3 (-0.967, 70000)
 -------------------------
```

Figure 11: Run `21-05-08_17:18:16` showing rewards for Huber `beta` being 1 (blue), 0.5 (orange) and 0.1 (pink). Smoothing 0.9.

Figure 12: See narrative on `21-05-08_18:30:55`. The lines are respectively orange, pink, blue, green. Smoothing 0.9.

marginally favouring noise of 0.1. See Figure 12.

6.2 Varying `ALICE_DOUBLE`, `N_NUMBERS` and `N_CODE`

Then ran at `21-05-08_23:17:59` with

```
'ALICE_DOUBLE': [None, 100, 300, 1000, 3000],
'N_CODE': [8, 16],
'N_NUMBERS': [16, 256]
```

getting

```
---- Table of results ----
```

```
      code hp_run        result
  00000000    1 (-0.998, 70000)
  00000001    2 (-0.994, 70000)
  00000010    3 (-0.979, 70000)
  00000011    4 (-0.984, 70000)
  00000100    5 (-0.992, 70000)
  00000101    6 (-0.993, 70000)
  00000110    7 (-0.963, 70000)
  00000111    8 (-0.977, 70000)
  00000200    9 (-0.994, 70000)
  00000201   10 (-0.994, 70000)
  00000210   11 (-0.907, 70000)
  00000211   12 (-0.972, 70000)
  00000300   13 (-0.902, 60000)
  00000301   14 (-0.963, 70000)
  00000310   15 (-0.654, 40000)
  00000311   16 (-0.847, 70000)
  00000400   17 (-0.810, 70000)
  00000401   18 (-0.909, 70000)
  00000410   19 (-0.644, 40000)
  00000411   20 (-0.728, 40000)
  ------------------------
```

ALICE_DOUBLE: None is the best, which deals pretty easily with all the sub–options, working marginally better with N_CODE: 8 for both N_NUMBERS: 16 and N_NUMBERS: 256. However, perhaps need to run longer than 70 000 iterations for higher N_NUMBERS and N_NUMBERS? Also, perhaps some ALICE_DOUBLE, say 100 to 1000, may help stabilise? Could revert to ALICE_DOUBLE: 500 as in 21-05-08_18:30:55. May also need to pay more attention to the Bob side of things.

For session 21-05-09_12:21:11 (terminated early), with N_CODE: 8, N_NUMBERS: 256 increasing N_SELECT: 16 to N_SELECT: 256 makes it a lot harder. More in the 0.7s and flat. It's also much slower has Bob has to try 256 rather than 16. However, 21-05-09_17:56:02 with four alternative tuples of seeds, does everything quickly and to near 1. It also suggest DOUBLE None or 500 doesn't make much difference — may None slightly more successful overall and should also be chosen on Occam principle.

Maybe set Python seed too?

7   Using MaxTempered layers

Defined MaxTempered layers, of which MaxTemperedInFocused seems better than MaxTemperedOutFocused, in src.lib.max_tempered_layers.py, and tried this out with Alice in 21-05-11_17:21:41 which seems as good as 21-05-08_23:17:59's first hp_run, albeit doubtless slower.

For N_SELECT=256, the runs 21-05-12_21:01:48 21-05-13_09:05:41 compare a feed forward and max layer networks— see Figure 13. So with 200 000 iterations all about the same. Recall that we switch off noise just over a bufferful before the end of the run.

Now with 70 000 at 21-05-13_14:17:04 compare feed forward with MaxLayer, now with relu=True, find the MaxLayer with or without bias_included does better than feed forward (essentially 1.0 rather

Figure 13: From `21-05-12_21:01:48`, MaxLayer with `bias_included=False` (red) and MaxLayer `bias_included=0.5` (grey) plus `21-05-13_09:05:41` feed forward (pink). Smoothing 0.9.


Figure 14: From `21-05-13_14:17:04`, feed forward `FFs(3, 50)` (red), `MaxNet("In", 3, 50, relu=True)` (grey) and `MaxNet("In", 3, 50, bias_included=True, relu=True)` (cyan). Smoothing 0.9.

than 0.95). However takes about two hours rather than one–and–a–half. Whether `MaxLayer` is better depends on how few iterations needed to get feed forward or `MaxLayer` to 1.0. Note that all the `MaxLayer` so far have `beta=0.2`, which means that its layers are 0.8 a normal feed forward are 0.2 a max–tempered layer — so should try higher beta too. See Figure 14.

Dropout instead of ReLU doesn't seem to help — see `21-05-14_22:23:16`.

So the best configuration appears to be that of `21-05-09_17:56:02`

```
{ #TODO enable dictionary-based choices for finer choosing
  'N_ITERATIONS': 70 * (10 ** 3), # 5 * (10 ** 5),
  'RANDOM_SEEDS': [
  (714844, 936892, 888616, 165835),
  (508585, 487266, 751926, 247136),
  (843402, 443788, 742412, 270619),
  (420915, 961830, 723900, 510954)
  ],
  'ALICE_NET': 'FFs(3, 50)',
  'BOB_LAYERS': 3,
  'BOB_WIDTH': 50,
  'BATCHSIZE': 32,
```

Figure 15: blue lower, green lower, cyan, orange, pink, blue, green, red

```
    'GAMESIZE': 32,
    'BUFFER_CAPACITY': 32 * 20000,
    'START_TRAINING': 20000,
    'N_SELECT': 256, #16,
    'EPSILON_ONE_END': 2000, #25000, # 40000,
    'EPSILON_MIN': 0.0,
    'EPSILON_MIN_POINT': 20000, #3 * (10 ** 5),
    'ALICE_PLAY': 'QPerCode',
    'ALICE_TRAIN': 'QPerCode', # 'FromDecisions',
    'BOB_STRATEGY': 'circular_vocab',
    'ALICE_OPTIMIZER': [
    'SGD(lr=0.01)'
    ],
    'BOB_OPTIMIZER': [
    ('SGD', '{"lr": 0.01}')
    ],
    'ALICE_LOSS_FUNCTION': [
    'Huber(beta=0.1)'
    ],
    'BOB_LOSS_FUNCTION': ('torch.nn.MSE', {}), # 'Same',
    'ALICE_PROXIMITY_BONUS': 10 ** 8, # 30000 * (10 ** 3),
    'ALICE_PROXIMITY_SLOPE_LENGTH': 10000,
    'ALICE_LAST_TRAINING': 100 * (10 ** 5),
    'NOISE_START': 30000,
    'NOISE': 0.1,
    'ALICE_DOUBLE': [None, 500],
    'N_CODE': 8,
    'N_NUMBERS': 256
}
```

which got

```
    ---- Table of results ----
```

```
   code hp_run result
   00000    1 (-0.986, 70000)
   00001    2 (-0.987, 70000)
   10000    3 (-0.978, 70000)
   10001    4 (-0.985, 70000)
   20000    5 (-0.984, 70000)
   20001    6 (-0.973, 70000)
   30000    7 (-0.990, 70000)
   30001    8 (-0.979, 70000)
```

Abandon `DOUBLE` on Occamist grounds.

Dropping noise seems to worsen to low 0.90s — see `21-05-15_07:57:53`.

For `N_CODE=4`, get `0.973` in the run at `21-05-15_11:49:36` at iteration 70 000, in two hours. At the end of this run Alice uses 13 codes.

For `N_CODE=2`, get `0.831` in the run at `21-05-15_14:20:40` at iteration 70 000, in two hours. At the end of this run Alice uses 4 codes.

For `N_CODE=8` again, with `21-05-15_17:16:53` get essentially 1 in 35 000 iterations. (Note forgot to shorten smoothing length so this result is evident in TensorBoard, not in the log.) Took 55 min.

With `21-05-15_17:16:53` get 0.975 in 17 500 iterations. So is beginning to degrade. Took 27 min.

With `21-05-15_20:18:01` get 0.982 in 25 000 iterations. Good enough? Took 39 min.

30 000 iterations gives 0.981. Stick with 25 000.

8   Things to try

1. What codes does best Alice generate?

2. Try using the loss function to constraint outputs to nearer bit values. Try increasing the weighting of this.

3. How quickly can `epsilon` be tapered?

4. Vary learning rates.

5. Vary `modulus`, `N_CODE` and `N_SELECT`.

6. Introduce noise.

7. Alice strategy with a code, as input and the output are values for the numbers. In each play (or train?) step feed all the codes in and the outputs indicate how well represents each number???

8. Try best strategy but with Alice outputs having dimension 2 ∗∗ `N_CODE`.

9. Train bits successively.

10. Look at MARL literature.

11. (At some stage in the training) introduce a 'proximity bonus' into Alice's training, which increases (in the same way) both the closeness of codes and the rewards if Bob's decision is right or nearly so.

12. Do a second sweep of `epsilon` going from high to low — perhaps for one player only? Definitely should re-`epsilon`–randomise Bob as otherwise Alice will never (or rarely if `N_SELECT < N_CODE`) get fed choices not in his decoding book. And I think Alice too, so Bob can learn new codes.

13. Random seeds seem to play a significant role — at least for short ($\sim 12\,500$) iteration training. Test how significant for $500\,000$ iterations.

14. Simulate use of a code–decode book pair.

References

[1] J. He, J. Chen, X. He, J. Gao, L. Li, L. Deng et al., *Deep reinforcement learning with a natural language action space*, *arXiv preprint arXiv:1511.04636* (2015) .

[2] G. Dulac-Arnold, R. Evans, H. van Hasselt, P. Sunehag, T. Lillicrap, J. Hunt et al., *Deep reinforcement learning in large discrete action spaces*, *arXiv preprint arXiv:1512.07679* (2015) .

[3] S. J. Majeed and M. Hutter, *Exact reduction of huge action spaces in general reinforcement learning*, *arXiv preprint arXiv:2012.10200* (2020) .